

Reproducibility by Other Means: Transparent Research Objects with Science-Oriented Provenance

Timothy McPhillips Bertram Ludäscher

School of Information Sciences, University of Illinois at Urbana-Champaign

{tmcphill,ludaesch}@illinois.edu

1 INTRODUCTION

Publicly funded efforts around the world currently are underway to ensure that computational components of scientific research can be made "reproducible" or "replicable". The ongoing discourse about the perceived "reproducibility crisis in science" is just one illustration of the importance of these efforts. The energy invested in the wide-ranging debate over the precise meanings of the terms *reproducible*, *replicable*, *transparent*, etc, with respect to research results, processes, and settings is perhaps an even greater indication of both the significance of these efforts and the challenges they face. For while each effort aimed at facilitating reproducible computing in the sciences must clearly define its mission and apply the bulk of its resources to the specific problems it sets out to address, these efforts necessarily do so within the context of broader discussions about the nature, importance, and precise definitions of the qualities of science we wish to extend to computing over the longer time scale.

Within a particular effort it is useful to define terms such as *reproducible* operationally. For example, in the Whole Tale project we define a *Reproducible Tale* as one that *includes sufficient information for the Tale to be re-executed for the review and verification of results*. Adopting this definition allow us to focus our requirements analysis, system design, and software implementation efforts on the specific problems Whole Tale is funded to solve and the use cases we aim to support. Supporting publishers who request authors to include all new data, code, and workflows needed to reproduce computed artifacts supporting claims in a paper is one such use case targeted by Whole Tale. We anticipate that facilitating re-execution of code used to generate key products of a study will enable publishers to routinely confirm that provided data and code do in fact produce those results—thereby addressing a key dimension of the reproducibility challenge currently facing science.

At the same time, it is critical that efforts like Whole Tale contribute to a global vision of computational reproducibility in the sciences, and clearly situate its particular mission, use cases, and engineering deliverables in this context. For while the particular technical problems that Whole Tale and similar projects aim to address are particularly pressing, current efforts by no means represent the entire landscape of concepts, problems, and technical options that will require further discussion, clarification, and analysis if we are to meet the challenges of reproducibility now facing the sciences. In particular, current engineering efforts are unlikely to elevate the computational components of research to the level of reproducibility expected of studies in the pure natural sciences such as physics, chemistry and biology.

Consequently, we view the Whole Tale project—as currently chartered and funded—as just a step towards the kind of platforms, infrastructure, and standards needed to enable researchers using

computing technology to routinely achieve the reproducibility long considered the essence of science as a whole. In support of this longer-term vision, we outline in this paper just a few of the issues we aim to investigate and discuss with the broader community over the next few years. We anticipate that future iterations of Whole Tale and its sibling efforts will be driven in part by the problem definitions and solution proposals we collectively develop between now and then.

2 OUTLINE OF PAPER

In the remainder of this paper we briefly discuss four topics we plan to investigate in the course of the Whole Tale project. In Section 3 we review the general notion of reproducibility in science and highlight how digital computing in principle makes possible a completely new kind of reproducibility: *exact repeatability*. We emphasize that the notion of *transparency*—long a critical element of reproducibility in the pure natural sciences—has a role to play even for those computational components of research where exact repeatability is feasible. In Section 4 we provide an overview of several dimensions of the terminological debate around reproducibility generally, and propose that a pluralistic approach to defining key terms is essential if a general concept of reproducibility is to be shared across disciplines. In Section 5 we summarize a number of limitations on exact repeatability in practice, and in Section 6 show how science-oriented provenance queries can mitigate such limitations by maintaining the transparency most essential to reproducibility in science. Throughout, we highlight the role that Research Objects can serve in supporting and maintaining reproducibility by encapsulating the information needed to rerun the computational steps in a study, and by enabling transparency via science-oriented queries of provenance information packaged in the object.

3 REPRODUCIBILITY IN SCIENCE

Modern science is founded on the expectation that the observations, experiments, and predictions that comprise scientific research be independently verifiable by others. This requirement, referred to as the *reproducibility* of science, applies not only to the products of research studies (*results, conclusions, models, data products, predictions*), but also to the activities that ultimately give rise to these products (*methods, protocols, workflows*); the materials employed in these activities (*reagents, instruments, software*); and the conditions under which these activities are carried out (*instrument settings, software parameters, computing environments*). When sufficient details are available such that the research products and methods can be reviewed, interpreted, and evaluated by other researchers *without* repeating the work, a study is said to be *transparent*.

While it is true that studies attempting primarily to reproduce previous results are relatively rare in the pure natural sciences, even the most groundbreaking studies in these fields include components that explicitly or implicitly confirm the reproducibility of previously reported results and procedures. The expectation is that new studies will reliably produce meaningful results consistent with previous work only if the prior work on which they are based or relates to is reproducible. In this sense, the whole of basic research in the natural sciences can be seen as an ongoing, massively-parallel reproducibility study that also happens to produce new results. Exceptions to this pattern occur when studies appear to overturn well-established understandings of nature, violate the expectations of how research in a particular field is to be carried out, or otherwise cause controversy. In these cases direct attempts may be made to reproduce results by precisely replicating the reported methods and conditions reported in the controversial study.

Even when attempts are made to confirm the reproducibility of particular studies or results, investigators in the natural sciences generally do not expect the processes and products of research to be repeated exactly. The vast majority of quantitative observations made of real world phenomena using scientific instruments are associated with limited precision and other intrinsic uncertainties that must themselves be characterized and well understood for science based on them to be considered reproducible. It is a hallmark of trustworthy science that quantitative observations and claims are inseparable from these uncertainties in measurement and their propagation through data analysis.

In contrast, digital computing approaches make it possible to repeat *exactly* certain computational aspects of research, far exceeding the precision achievable when observing natural phenomena in the physical world. It generally is expected that computational processes, the implementation of hardware and software enabling those processes, and the outputs of those processes all can be repeated exactly—at least in principle. Where computing plays a significant role in scientific research, there is risk of this new expectation of exact repeatability being conflated with the longstanding understanding of reproducibility in the basic sciences.

Moreover, while computational experiments and analyses may be exactly repeatable in principle, in practice the complexities of real-world hardware and software make computational repeatability challenging to achieve in practice except over very limited time scales. Because of the obvious value that exact repeatability brings when it is feasible, it is important that we work to expand the fraction of scenarios in which the computational components of research can be automatically repeated exactly over ranges of time and space relevant to scientific research and discourse. These efforts are particularly important for the research community to pursue, and for science funding agencies to support, because the computing industry generally does not have requirements for exact repeatability across significant spans of time. It is important to note, however, that the concept of exact repeatability is qualitatively different from the concept of reproducibility that underlies the natural sciences. In particular, scientific reproducibility is not simply a weaker form computational repeatability. Achieving or approximating computational repeatability does not automatically deliver scientific reproducibility.

It is in a sense both bad and good news that exact computational repeatability is not tantamount to reproducibility. The disappointing news, perhaps, is that it is possible to put much effort into achieving computational repeatability, exact or even inexact, without delivering the kind of reproducibility that is critical for producing trustworthy science. The good news is that scientifically-meaningful reproducibility can be realized in cases (or over spans of time) where computational repeatability is either impossible or impractical due to limitations on available technology and affordable resources. In this paper we propose that the older concept of reproducibility that permeates the basic natural sciences has a very useful role even where digital computing makes exact repeatability a theoretical possibility.

Researchers in the natural sciences are comfortable with the idea that it is not possible to exactly repeat all reported observations, procedures, and experimental results. They do not see this as a contradiction to their demand that science be reproducible. What the natural sciences actually do demand is that (a) research procedures be repeatable by others in principle; (b) the means of repeating the work be subject to review and evaluation; and (c) such review and evaluation be possible *without* actually repeating the work. To be clear, in the natural sciences it is actually considered a *problem* if exact repetition of the steps taken in reported research is required either to evaluate the work or to reproduce results.

Consequently, it is not necessary to achieve or maintain perfect repeatability of the computational components of research for scientists to consider a study reproducible and therefore trustworthy. At that same time it is important that the standards, technologies, computational best-practices, and infrastructure we develop and advocate in fact support scientific reproducibility. It is not enough to pursue and support exact computational repeatability where we can, and to get as close as possible to exact reproducibility otherwise. Rather, computational repeatability is best seen as a dimension of research reproducibility *orthogonal* to the dimension of transparency. It is possible to achieve computational repeatability without providing research transparency, and vice versa. And while exact repeatability is not an essential element of reproducibility, transparency is.

In this paper we argue that the dimension of reproducibility most ripe for the contributions of computer science is research *transparency*, in particular through the modeling, recording, and querying of the provenance of research artifacts. In alignment with researchers in the natural sciences who recognize transparency as crucial, we are confident that provenance management has much to contribute to scientific reproducibility, even when it does not specifically enable exact repeatability of the computations they describe.

For provenance management systems, representations, and user interfaces to support reproducibility via transparency, however, they must support science-oriented queries. Provenance must be able to answer questions about the *science* that was performed—not just the sequence, dependencies, and flow of data through computational steps. The answers to these questions must enable others to evaluate the scientific quality of the work, and to learn what is necessary to reproduce the results without actually repeating every step taken in the original work. Provenance must enable

researchers to build on the results and processes reported in prior work with confidence.

Finally, it must be possible for researchers unversed in the detailed specifications of Research Objects and the PROV standard to pose questions and receive answers meaningful for evaluating, using, and building on the processes and products of prior research. We suggest that Research Objects and related approaches are the ideal vehicle for storing, sharing and making provenance queryable in this way. Research Objects thus can support scientific reproducibility even in the face of the many practical challenges to computational repeatability.

4 TERMINOLOGY

Despite broad acceptance of the importance of general idea of reproducibility sciences, recent years have seen a growing debate over the precise meanings of the terms *reproducible* and *replicable*, and the relationship between the two. For example, should one term imply a greater adherence to the procedures and methods employed in the original research? Should one term be reserved for the activities a researcher performs to assess the reliability of their own procedures, experiments, and results; leaving the other term to refer to activities centered on assessing the quality of the results of others? Which term should be used to mean what? And what value is there in distinguishing reproducibility from replicability at all?

As an example of an effort to distinguish reproducibility and replicability the FASEB recommendations are illuminating:

Replicability: the ability to duplicate (i.e., repeat) a prior result using the same source materials and methodologies. This term should only be used when referring to repeating the results of a specific experiment rather than an entire study.

Reproducibility: the ability to achieve similar or nearly identical results using comparable materials and methodologies. This term may be used when specific findings from a study are obtained by an independent group of researchers

Footnote: The Federation of American Societies for Experimental Biology comprises twenty-nine distinct scientific societies together representing over 105,000 practicing researchers.

According to these definitions, *replicability* indicates a higher level of fidelity than does *reproducibility*, both with respect to the prior result to be confirmed, and to the materials and methodologies employed. Replicability also appears more likely for the original researchers (they presumably have access to the "same source materials" and are in the best position to use the same methodologies), whereas reproducibility is considered feasible for "an independent group of researchers".

Neither definition refers specifically to computed results and both apply to experimental results. In contrast, the recent report from the National Academy of Sciences reverses the relative fidelity implied by the terms 'reproducibility' and 'replicability':

Reproducibility is obtaining consistent results using the same input data, computational steps, methods, and code, and conditions of analysis.

Replicability is obtaining consistent results across studies aimed at answering the same scientific question, each of which has obtained its own data.

Perhaps most intriguing about the NAS definitions is that experiments not carried out entirely in silico apparently are left with only term to describe them: replicability. Satisfying the definition of reproducibility requires "computational steps" and "code", and the report goes on to clarify that reproducibility "is synonymous with 'computational reproducibility,' and the terms are used interchangeably in this report."

FASEB also defines transparency in a manner consistent with our usage here:

Transparency: the reporting of experimental materials and methods in a manner that provides enough information for others to independently assess and/or reproduce experimental findings

In contrast, the NAS report essentially makes reproducibility a prerequisite of transparency: "When a researcher transparently reports a study and makes available the underlying digital artifacts, such as data and code, the results should be computationally reproducible."

We take the position that it transparency is essential to science, whereas reproducibility as defined by the NAS is not necessarily possible, as detailed below.

Footnote or note at end of paper:

Cell and molecular biologists study both replication and reproduction as natural processes. DNA replicates : high fidelity, variation not desired, ingredients indistinguishable, errors are corrected on the fly. Organisms reproduce: lower fidelity expected, variation desired, different ingredients acceptable. Cells have replisomes, complex molecular machines where DNA replication occurs, and copying errors are detected and corrected. In origin of life research a crucial debate is over 'replication first' (DNA World) or 'metabolism first' (aka reproduction first, i.e. without replicating genetic material). These terminologies with biology are well established. Biologists also have a rich and well-defined vocabulary to describe replicability of experimental measurements and results. Commonly distinguish two kinds of experimental 'replicates': technical replicates, and biological replicates. Reality is that the terminology is well established in large branches of science already.

Need for reproduction/replication to mean different things in different fields

The relationships of corresponding concepts across fields is one of analogy, not identity. Exact repeatability is at least theoretically possible and sometimes practical under realistic assumptions when an experiment is entirely in silico and isolated from the outside world. As soon as observation of the real world is involved, exact repeatability often is impossible. Neither of the types of replicates in experimental biology are exact, although both are measures of repeatability. There often is no way to repeat exactly an experiment that involves scientific instruments, physical samples, or experimental apparatus. In contrast, it is not unreasonable to talk about exactly repeating a purely computational experiment, at least by the original researcher, on the same hardware, close in time to the original experiment. In reality, computational repeatability is not as easy as sometimes assumed (see below), but fundamentally this is a different situation than when a scientific instrument is involved or observations are made of the external world.

Our approach to terminology

Respect differences between fields of research and different expectations with regard to reproducibility. Do not expect standardization to even be meaningful (never mind politically achievable) across domains. Instead, only use the R* words in ways that make their meanings clear in context. Do not be surprised if computational sciences turn out not to be representative of science generally.

Specific implications for our contributions to the Research Objects field

Take care to define R* words precisely when expressing desiderata, describing features, or making or comparing claims about capabilities. Do not expect efforts to achieve computational repeatability alone to enable "reproducible science" generally.

As alluded to in the introduction, the entry of digital computing into scientific research has further complicated these issues somewhat by introducing the possibility of exact repeatability. For the purposes of this paper we will continue to reserve the terms reproducibility and replicability to qualities demanded of scientific research whether digital computers are involved or not. We use the term repeatability to refer to quality of computations that is desirable so far as computers are involved in research, but do not consider repeatability in this sense to be a prerequisite of reproducibility or replicability even when computers are used in research.

5 COMPUTATIONAL CHALLENGES

The fundamental limitations computers impose on replicability of program executions are well known.

Finite precision arithmetic, different word sizes on different processors, round-off errors, etc, impose limits on scientific computations and their replicability across different computing environments. Virtual machines and containers do not address these issues. Full emulation is required to run the same binary in identical fashion on a different processor. This is typically slow. These limitations are even more challenging to manage reproducibly because programs typically are compiled, meaning that the exact sequence of machine instructions executed even by a single processor cannot generally be controlled. A different compiler, or a newer version of the same compiler will yield a different sequence of machine instructions.

Replicating the outputs of a program is far from straightforward

Observing that a program or set of programs can be executed again is not sufficient to conclude that the underlying computation was replicated. The outputs of the programs must be checked for equivalence. Because of the expected variation in run time behavior of programs due to the issues above, checking that outputs of a program run are equivalent to the outputs of a run of that program is not always as simple as comparing the outputs for bitwise identity. Robustly checking for equivalence of output generally must be confirmed in some way other than comparing files at the bit level. Footnote: The excellent practice of including accurate provenance and other meaningful metadata in data file headers makes it even more unlikely that outputs from different runs will be bitwise-identical.

Replicating just the software running the program is challenging in practice

How can we ensure that the stream of instructions sent to the processor for two executions is identical? Even holding the computer hardware and compiler version constant, programs depend on language libraries, OS libraries, and system calls. Much scientific software also depends directly and indirectly on large numbers of 3rd-party libraries. Only direct dependencies can be controlled reliably at build-time. And many dependencies are via shared libraries that can change between executions of the exact same executable—no recompile is needed to get a new effective executable. Footnote: Fans of the Go programming language are bringing back the static executable. Recompiling or even just rerunning the "same" program a week later can result in a completely different effective instruction stream.

Even reproducing computing environments is hard

Containers and their discontents Footnote: By 'discontents' we do not mean that we object to the use of containers, but rather than we are not content with container technology alone. There currently is much enthusiasm around containers as a means of reproducing computing environments and making computational science replicable. Whole Tale is one of several projects leveraging the capabilities of containers for this purpose. Others include Binder and Code Ocean. In Whole Tale it is recognized that containers alone cannot satisfy researcher's needs for sharing their computing environments and computations. Rather, container technology such as Docker provide an invaluable tool for the reproducible science software stack architect. A major motivation for funding (and continuing to fund) projects like Whole Tale is that the containers on their own are insufficient as means to making computational science reproducible, and it is not practical for individual researchers and groups to use containers and other technology to actually achieve scientifically meaningful reproducibility over periods longer than the publication-cycle time scale.

What containers do not do Despite what sound like suggestions to the contrary in the literature, container technology such as Docker do not ensure computational replicability, and do not on their own solve any of the problems of computational replicability described above. What containers do provide a very convenient means for executing customized computing environments on behalf of researchers, without having to run an entire virtual machine for each environment. In common with virtual machine technology, containers do not abstract or hide the underlying hardware architecture of the computer on which they run. They do not abstract the underlying operating system, but simply use the Linux kernel on the host. Kernel parameter settings on the host apply to all containers running on the host (reference famous blog post on the topic "Containers Do Not Contain"). Rebuilding an image from its Dockerfile specification is not guaranteed to yield the same image. In general it will not. Container images, once they are built, are not guaranteed to run on future releases of the container host. They also do not ensure that computations run within the container will be replicable in the future.

What containers are for What containers are good for is precisely what they were to do for the computing industry: enable developers to write and test code in a computing environment of a developer's choosing that can then be replicated on a very short time scale (hours or days) in staging and production environments. Containers also are good at managing conflicts in dependencies

Reproducibility by Other Means

between different components of a multiprocess software architecture. Using containers to 'contain' dependencies in this way is most effective when an application can be split across multiple containers running in concert. The model of one container, one computing environment does not lend itself to dependency isolation.

The problem of time and dependencies An emerging threat to reproducibility of computational science is the spread misconception that sharing the definition of a container image, e.g. by including a Dockerfile in the Git repo for the project, is a guarantee that others (or even the original researcher) will be able to recreate the corresponding image and computing environment it represents. Researchers making this assumption may be less likely to preserve all of the information actually required to reproduce their computations. A major reason a Dockerfile is not enough is that the implicit dependencies of the built environment are constantly changing.

This is well known to anyone working directly with Docker, or other container technologies. Here we will give a single example of the implications of this issue for reproducible science.

What is reproducibility really for?

Achieving meaningful replicability even for the computational parts of research is very challenging. But this is no reason to give up hope. Replicability is a means to an end—justification of scientific results—and Research Objects can help us achieve that end by other means. What is most exciting about Research Objects is that they can achieve this end despite the difficulty of computational replicability. And the primary means by which Research Objects can do this is by providing transparency.

REFERENCES